

Graph Algorithms (continued)

COMP4500/7500

Advanced Algorithms & Data Structures

August 21, 2019

Overview

- Admin/reminders
- Shortest paths:
 - Dijkstra's algorithm
 - Bellman-Ford algorithm
- Priority-first search

Shortest paths

A path p from v_1 to v_n is a sequence of vertices.

$$v_1 \rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_n$$

The total weight of the path, $\text{weight}(p)$, is the sum of the edges along the way.

$$\text{weight}(p) = \sum_{i=1}^{n-1} \text{weight}(v_i, v_{i+1})$$

The **distance** from v_1 to v_n is the minimum weight of all paths from v_1 to v_n .

A **shortest path** from v_1 to v_n is a path from v_1 to v_n of minimum weight.

Shortest paths

Types of shortest paths:

- ① **Single-pair.** Given a pair (u, v) , find a shortest path between them.
- ② **Single-source.** Given a vertex v , find a shortest path to every other vertex. *This is our focus.*
- ③ **Single-destination.** Given a vertex v , find a shortest path from every other vertex.
- ④ **All-pairs.** Given a graph, find a shortest path between every pair of vertices.

SINGLE SOURCE SHORTEST PATHS(G, w, s): key ideas

Given graph G with weight-function w and source vertex s ,
for each vertex $v \in G.V$ we calculate:

$v.d$ its **distance** from source vertex s

$v.\pi$ the **predecessor** to v on a shortest path from s to v

Steps:

- **Initialise** $s.d = 0$ and $v.d = \infty$ for all $v \in G.V - \{s\}$.
 - invariant: $distance(s, v) \leq v.d \leq \infty$
- **Initialise** $v.\pi = NIL$ for all $v \in G.V$.
 - invariant: $v.\pi$ is v 's predecessor on the shortest path found so far
- ... we **relax** our estimates until they reach a solution ...

Edge relaxation: relax edge $(u, v) \in G.E$

RELAX(u, v, w):

updates $v.d$ and $v.\pi$ if v can be reached faster via intermediate vertex u .

RELAX(u, v, w)

- 1 **if** $v.d > u.d + w(u, v)$
- 2 $v.d = u.d + w(u, v)$
- 3 $v.\pi = u$

- preserves the distance and predecessor invariants:
 - invariant: $distance(s, v) \leq v.d \leq \infty$
 - invariant: $v.\pi$ is v 's predecessor on the shortest path found so far
- does not change $v.d$ and $v.\pi$ if the shortest path to v has already been found

Edge relaxation

Let

$$p_n = v_1 \rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_n$$

be a shortest path from v_1 to v_n .

If edge (v_{i-1}, v_i) is relaxed *after* a shortest path to v_{i-1} is found, we will find a shortest path to v_i .

Graphs without negative weight edges

Let

$$p_n = v_1 \rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_n$$

be any **shortest path** from v_1 to v_n in a weighted graph with **no negative weight edges**.

For all $v_i \in v_1, v_2, \dots, v_{n-1}$, we must have that the path prefix

$$p_i = v_1 \rightarrow v_2 \rightarrow v_3 \rightarrow \dots \rightarrow v_i$$

is a **shortest path** from v_1 to v_i , and $\text{weight}(p_i) \leq \text{weight}(p_n)$.

i.e. $\text{distance}(v_1, v_i) \leq \text{distance}(v_1, v_n)$.

Relaxation for graphs without negative weight edges

If we:

- **relax** each edge $(u, v) \in G.E$ once
 - in order of u 's distance from start vertex s ,
- then we will find all the shortest paths from s in G .

DIJKSTRA(G, w, s)

Solves the **single-source shortest paths problem** for weighted graphs **without negative weight edges**.

[AI in computer games use a sophisticated variant called A^* .]

DIJKSTRA(G, w, s): key idea

For each vertex $v \in G.V$ we calculate:

$v.d$ its **distance** from source vertex s

$v.\pi$ the **predecessor** to v on the shortest path from s to v

Steps:

- **Initialise** $s.d = 0$ and $v.d = \infty$ for all $v \in G.V - \{s\}$.
 - invariant: $distance(s, v) \leq v.d \leq \infty$
- **Initialise** $v.\pi = NIL$ for all $v \in G.V$.
 - invariant: $v.\pi$ is a predecessor on shortest path found so far
- **Visit** each vertex in order of its distance from the source vertex – when we visit a vertex we **relax** its outgoing edges.

Generalises breadth-first search for weighted graphs.

DIJKSTRA(G, w, s)

How do we (efficiently) find the next vertex to visit?

Let

S the set of **visited vertices**

We maintain

Q a **priority queue** containing vertices $V - S$, with **key** $v.d$

At each step we visit the vertex u on Q with the minimum key.

When we visit vertex u , we are guaranteed to have already found a shortest path to u – *assuming that the graph has no negative weight edges!*

DIJKSTRA(G, w, s): why does the idea work?

When we visit a vertex u , we have already found a shortest path to u .

Let v be the predecessor of u on a shortest path from s to u .

$distance(s, v) \leq distance(s, u)$, and so either

- we have already visited v and relaxed its edges, or
- we already found another shortest-path to u of length
 $distance(s, u) = distance(s, v)$

Dijkstra's algorithm

DIJKSTRA(G, w, s)

```

1 //  $G$  is the graph,  $w$  the weight function,  $s$  the source vertex
2 INIT-SINGLE-SOURCE( $G, s$ )
3  $S = \emptyset$            //  $S$  is the set of visited vertices
4  $Q = G.V$            //  $Q$  is a priority queue maintaining  $G.V - S$ 
5 while  $Q \neq \emptyset$ 
6      $u = \text{EXTRACT-MIN}(Q)$ 
7      $S = S \cup \{u\}$ 
8     for each vertex  $v \in G.\text{Adj}[u]$ 
9         RELAX( $u, v, w$ )

```

INIT-SINGLE-SOURCE(G, s)

```

1 for each vertex  $v \in G.V$ 
2      $v.d = \infty$ 
3      $v.\pi = \text{NIL}$ 
4      $s.d = 0$ 

```

RELAX(u, v, w)

```

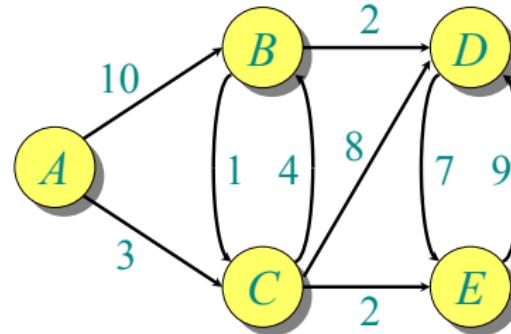
1 if  $v.d > u.d + w(u, v)$ 
2      $v.d = u.d + w(u, v)$ 
3      $v.\pi = u$ 

```



Example of Dijkstra's algorithm

Graph with
nonnegative
edge weights:

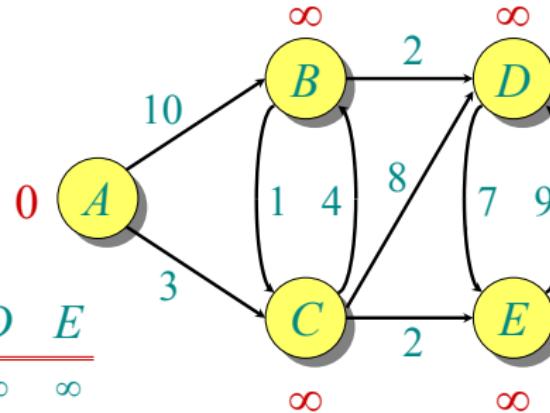




Example of Dijkstra's algorithm

Initialize:

$$Q: \begin{array}{ccccc} A & B & C & D & E \\ \hline 0 & \infty & \infty & \infty & \infty \end{array}$$

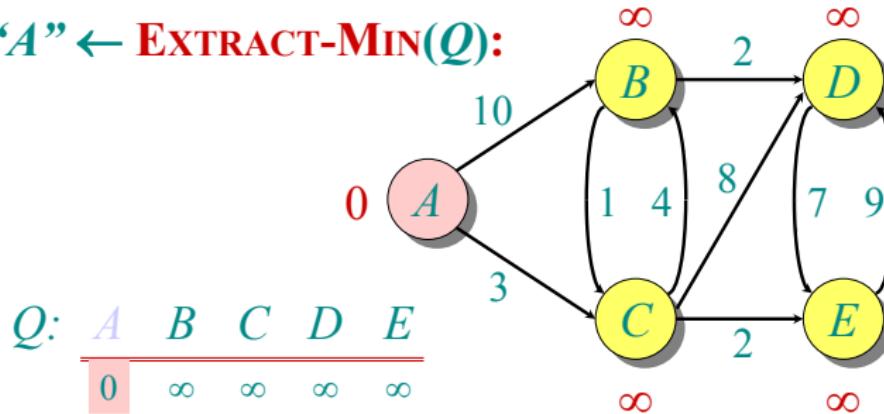


$$S: \{\}$$



Example of Dijkstra's algorithm

$"A" \leftarrow \text{EXTRACT-MIN}(Q)$:

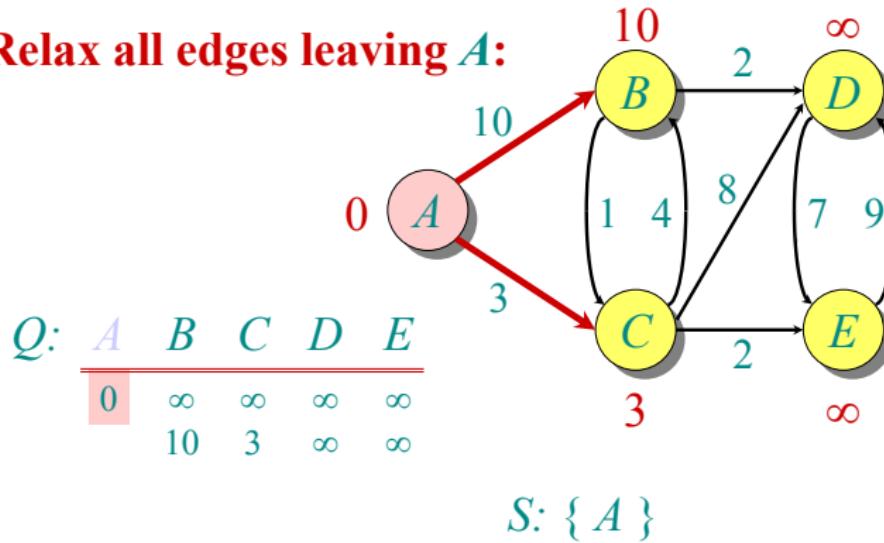


$$S: \{ A \}$$



Example of Dijkstra's algorithm

Relax all edges leaving A :

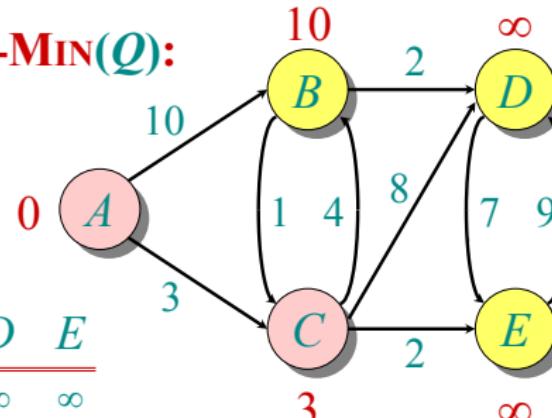




Example of Dijkstra's algorithm

$\text{“C”} \leftarrow \text{EXTRACT-MIN}(Q)$:

$Q:$	A	B	C	D	E
	0	∞	∞	∞	∞

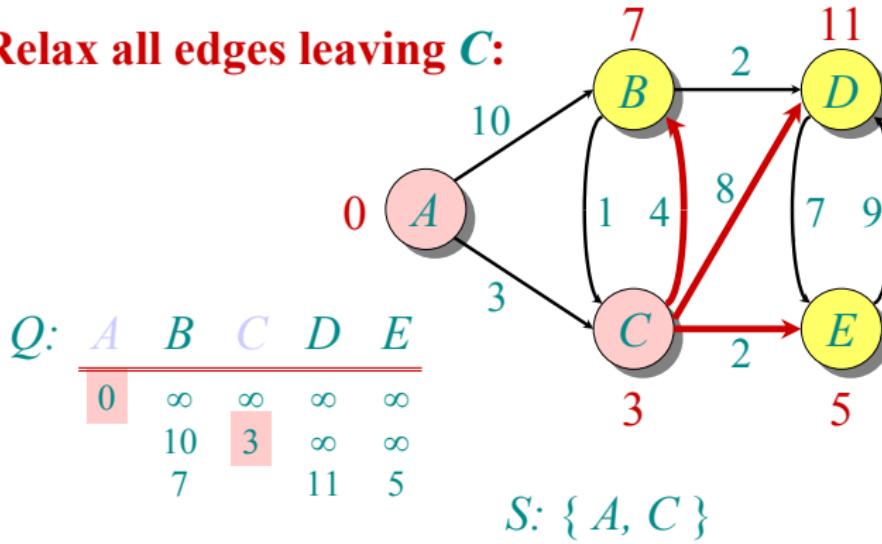


$S: \{ A, C \}$



Example of Dijkstra's algorithm

Relax all edges leaving C :

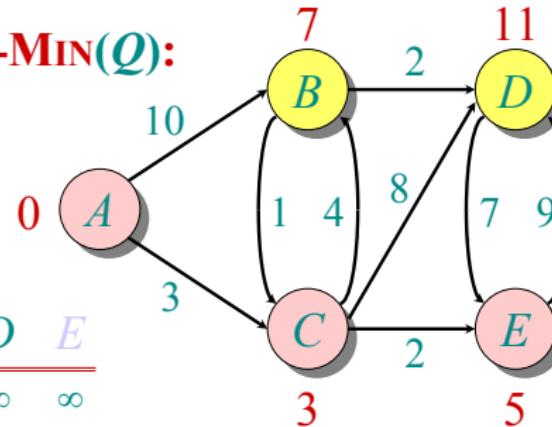




Example of Dijkstra's algorithm

$E \leftarrow \text{EXTRACT-MIN}(Q)$:

	A	B	C	D	E
A	0	∞	∞	∞	∞
B	10	3	∞	∞	∞
C	7		11	5	

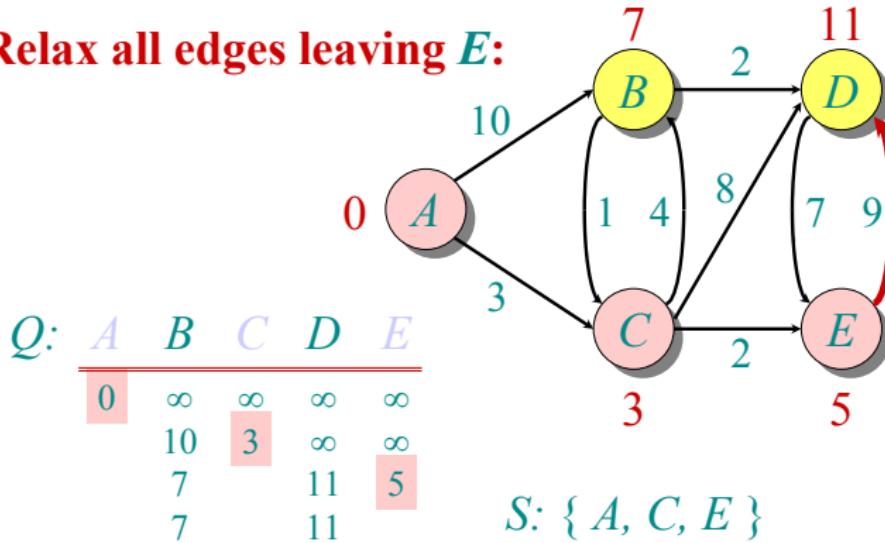


$S: \{ A, C, E \}$



Example of Dijkstra's algorithm

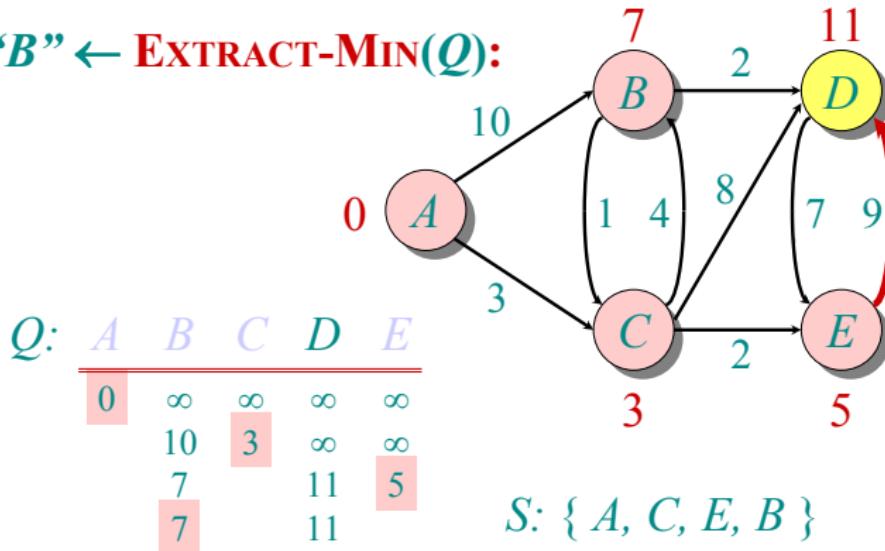
Relax all edges leaving E :





Example of Dijkstra's algorithm

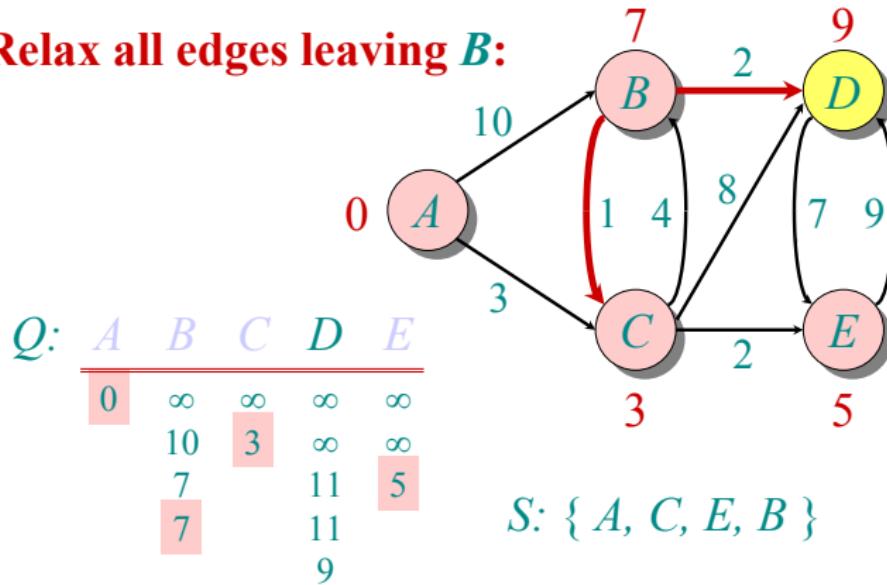
$\text{“B”} \leftarrow \text{EXTRACT-MIN}(Q)$:





Example of Dijkstra's algorithm

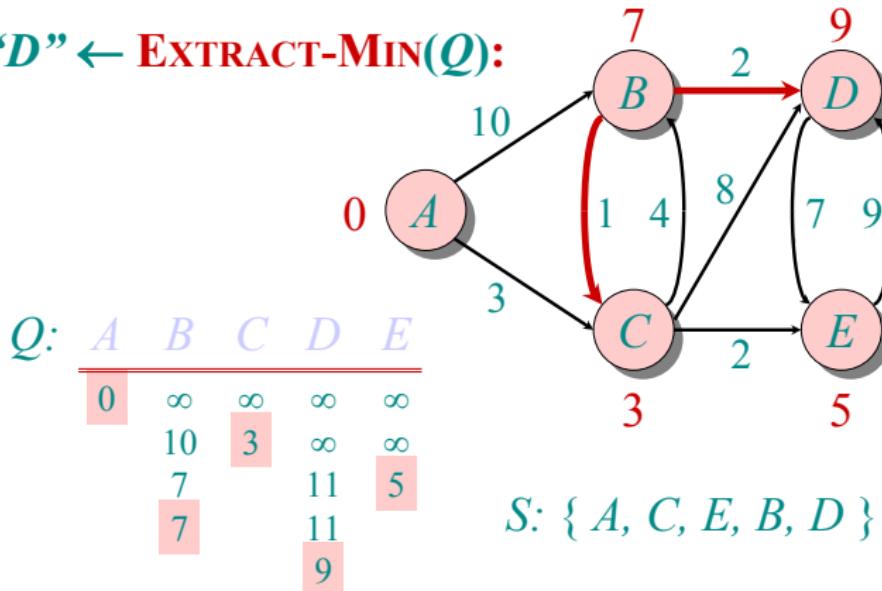
Relax all edges leaving B :

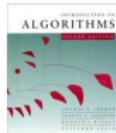




Example of Dijkstra's algorithm

$\mathcal{D} \leftarrow \text{EXTRACT-MIN}(Q)$:





Analysis of Dijkstra

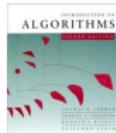
```
while  $Q \neq \emptyset$ 
    do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
         $S \leftarrow S \cup \{u\}$ 
        for each  $v \in \text{Adj}[u]$ 
            do if  $d[v] > d[u] + w(u, v)$ 
                then  $d[v] \leftarrow d[u] + w(u, v)$ 
```



Analysis of Dijkstra

$|V|$ times {

```
while  $Q \neq \emptyset$ 
    do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
         $S \leftarrow S \cup \{u\}$ 
        for each  $v \in \text{Adj}[u]$ 
            do if  $d[v] > d[u] + w(u, v)$ 
                then  $d[v] \leftarrow d[u] + w(u, v)$ 
```



Analysis of Dijkstra

```

while  $Q \neq \emptyset$ 
  do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
       $S \leftarrow S \cup \{u\}$ 
      for each  $v \in \text{Adj}[u]$ 
        do if  $d[v] > d[u] + w(u, v)$ 
            then  $d[v] \leftarrow d[u] + w(u, v)$ 

```

$|V|$ times $\left. \begin{array}{c} \\ \\ \end{array} \right\} \text{degree}(u) \text{ times}$



Analysis of Dijkstra

```

while  $Q \neq \emptyset$ 
  do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
       $S \leftarrow S \cup \{u\}$ 
      for each  $v \in \text{Adj}[u]$ 
        do if  $d[v] > d[u] + w(u, v)$ 
            then  $d[v] \leftarrow d[u] + w(u, v)$ 

```

$|V|$ times $\left\{ \begin{array}{l} \\ \text{degree}(u) \\ \text{times} \end{array} \right\}$

Handshaking Lemma $\Rightarrow \Theta(E)$ implicit DECREASE-KEY's.



Analysis of Dijkstra

```

while  $Q \neq \emptyset$ 
  do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
       $S \leftarrow S \cup \{u\}$ 
      for each  $v \in \text{Adj}[u]$ 
        do if  $d[v] > d[u] + w(u, v)$ 
            then  $d[v] \leftarrow d[u] + w(u, v)$ 
    
```

$|V|$ times {
 $\text{degree}(u)$ times {

Handshaking Lemma $\Rightarrow \Theta(E)$ implicit DECREASE-KEY's.

Time = $\Theta(V \cdot T_{\text{EXTRACT-MIN}} + E \cdot T_{\text{DECREASE-KEY}})$

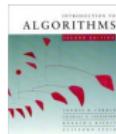
Note: Same formula as in the analysis of Prim's minimum spanning tree algorithm.



Analysis of Dijkstra (continued)

Time = $\Theta(V) \cdot T_{\text{EXTRACT-MIN}} + \Theta(E) \cdot T_{\text{DECREASE-KEY}}$

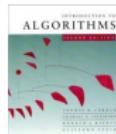
Q	$T_{\text{EXTRACT-MIN}}$	$T_{\text{DECREASE-KEY}}$	Total
-----	--------------------------	---------------------------	-------



Analysis of Dijkstra (continued)

Time = $\Theta(V) \cdot T_{\text{EXTRACT-MIN}} + \Theta(E) \cdot T_{\text{DECREASE-KEY}}$

Q	$T_{\text{EXTRACT-MIN}}$	$T_{\text{DECREASE-KEY}}$	Total
array	$O(V)$	$O(1)$	$O(V^2)$



Analysis of Dijkstra (continued)

Time = $\Theta(V) \cdot T_{\text{EXTRACT-MIN}} + \Theta(E) \cdot T_{\text{DECREASE-KEY}}$

Q	$T_{\text{EXTRACT-MIN}}$	$T_{\text{DECREASE-KEY}}$	Total
array	$O(V)$	$O(1)$	$O(V^2)$
binary heap	$O(\lg V)$	$O(\lg V)$	$O(E \lg V)$



Analysis of Dijkstra (continued)

Time = $\Theta(V) \cdot T_{\text{EXTRACT-MIN}} + \Theta(E) \cdot T_{\text{DECREASE-KEY}}$

Q	$T_{\text{EXTRACT-MIN}}$	$T_{\text{DECREASE-KEY}}$	Total
array	$O(V)$	$O(1)$	$O(V^2)$
binary heap	$O(\lg V)$	$O(\lg V)$	$O(E \lg V)$
Fibonacci heap	$O(\lg V)$ amortized	$O(1)$ amortized	$O(E + V \lg V)$ worst case

Negative-weight problems

What if our graph has negative weights?

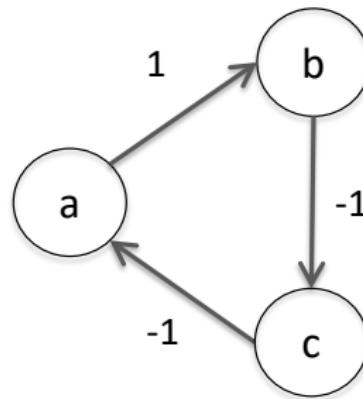
Dijkstra's algorithm might not relax edges in the right order.

I.e. (u, v) might be relaxed before the shortest path to u has been found.

Negative-weight problems

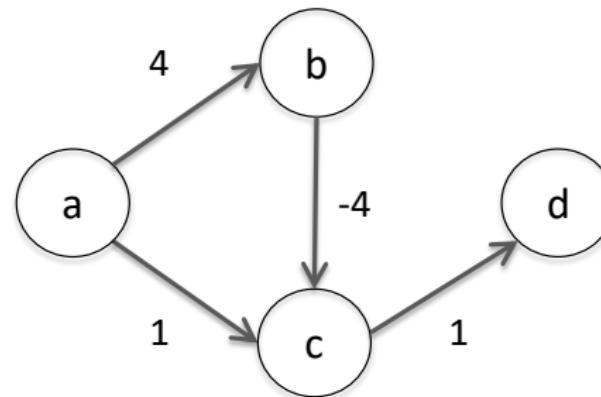
Are shortest paths well-defined for graphs with negative weight cycles?

What is the shortest path from vertex a to a ?



Quick question

Give a directed, weighted graph **with negative weights**, but no negative weight cycles, for which Dijkstra's algorithm fails.



Negative weight problems

- Graphs with negative weights might have **negative weight cycles**.
- Even if a graph with negative weights has no negative-weight cycles, we need to work out how to relax edges so that we are guaranteed to find shortest paths.

Negative weight problems: the Bellman-Ford solution

What do we know?

After initialization:

we have found all shortest paths that contain ≤ 0 edges.

After relaxing each edge in the graph 1 time:

we have found all shortest paths that contain ≤ 1 edges.

After relaxing each edge in the graph 2 times:

we have found all shortest paths that contain ≤ 2 edges.

...

After relaxing each edge in the graph $V-1$ times:

we have found all shortest paths that contain $\leq V-1$ edges.

Bellman-Ford algorithm

The Bellman-Ford algorithm uses this idea to find:

- single-source shortest paths on directed graphs
- that may have negative-weight edges.

Moreover, it

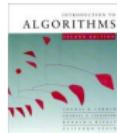
- detects negative weight cycles, returning *false* if one is found.

Bellman-Ford algorithm

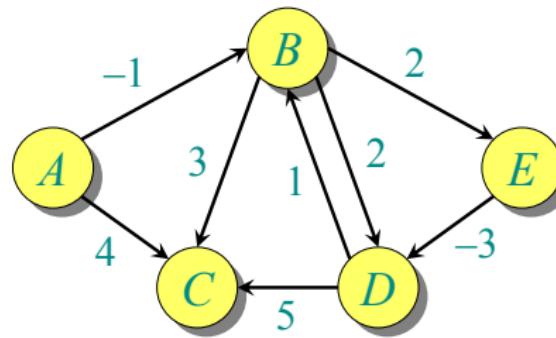
BELLMAN-FORD(G, w, s)

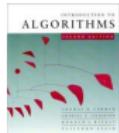
```
1 //  $G$  is the graph,  $w$  the weight function,  $s$  the source vertex
2 INIT-SINGLE-SOURCE( $G, s$ )
3 // Relax each edge  $V - 1$  times to find shortest paths
4 for  $i = 1$  to  $|G.V| - 1$ 
5     for each edge  $(u, v) \in G.E$ 
6         RELAX( $u, v, w$ )
7 // Check for negative-weight cycles
8 for each edge  $(u, v) \in G.E$ 
9     if  $v.d > u.d + w(u, v)$ 
10        return FALSE
11 return TRUE
```

What is the time complexity? $O(VE)$ (i.e., could be $O(V^3)$).

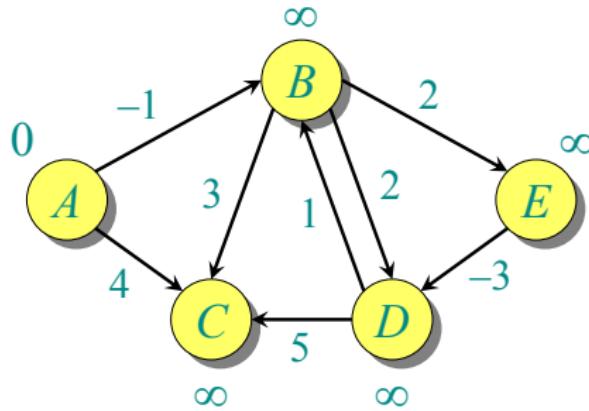


Example of Bellman-Ford

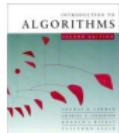




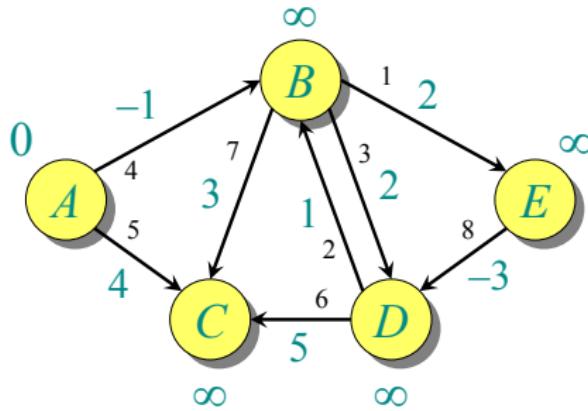
Example of Bellman-Ford



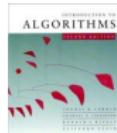
Initialization.



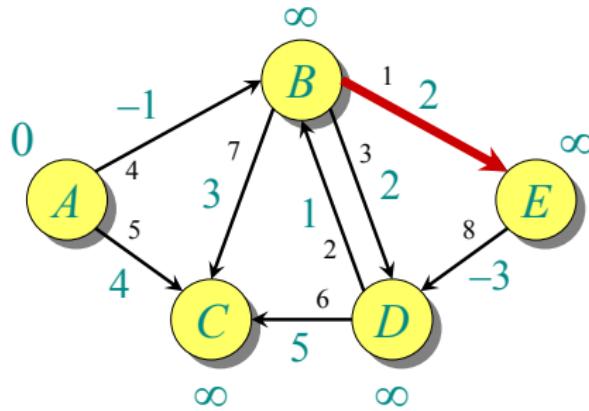
Example of Bellman-Ford

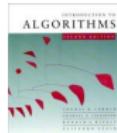


Order of edge relaxation.

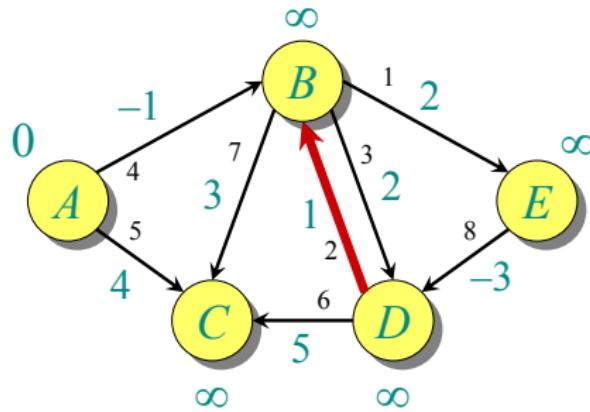


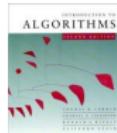
Example of Bellman-Ford



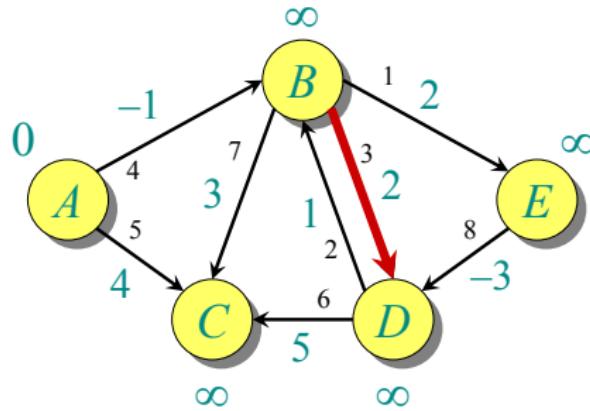


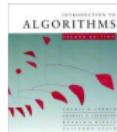
Example of Bellman-Ford



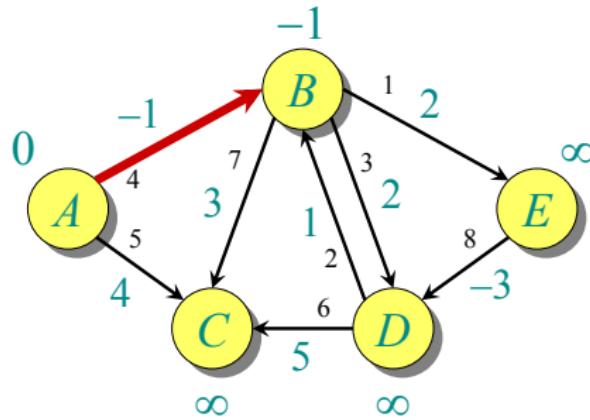


Example of Bellman-Ford



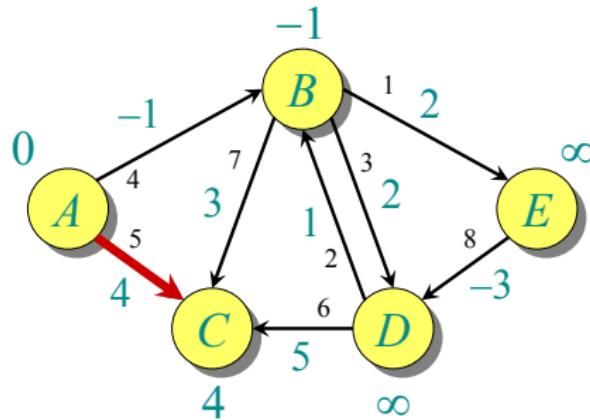


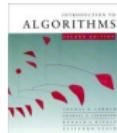
Example of Bellman-Ford



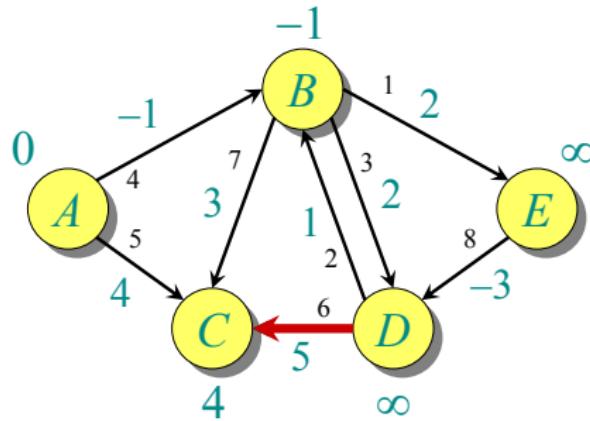


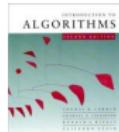
Example of Bellman-Ford



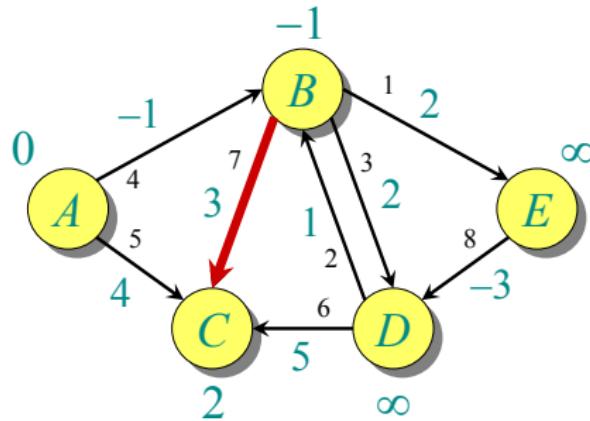


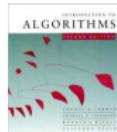
Example of Bellman-Ford



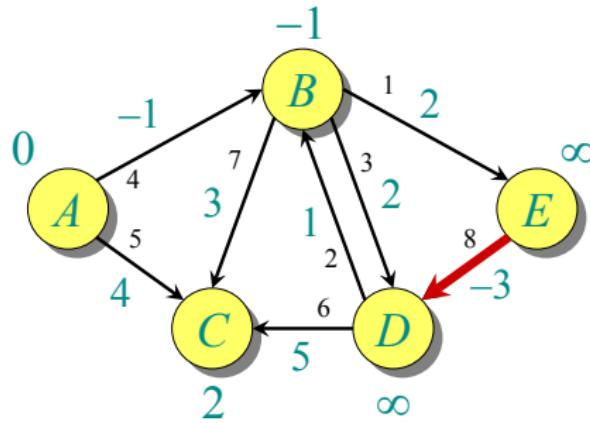


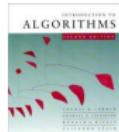
Example of Bellman-Ford



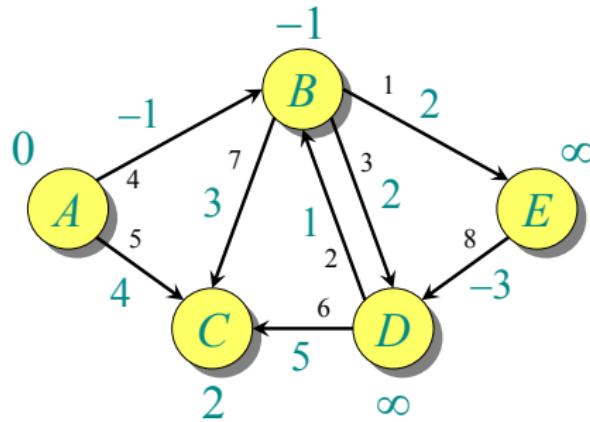


Example of Bellman-Ford

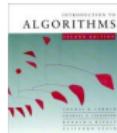




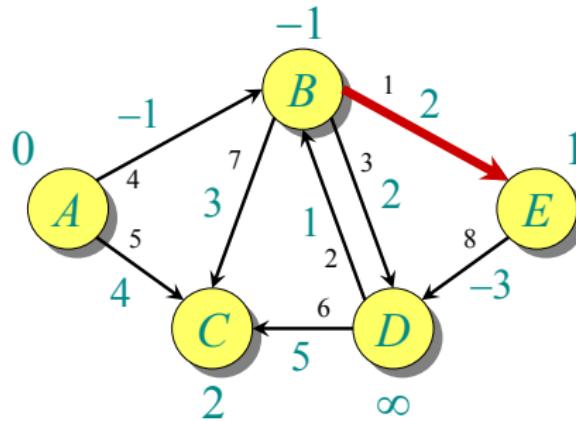
Example of Bellman-Ford

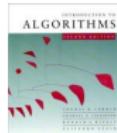


End of pass 1.

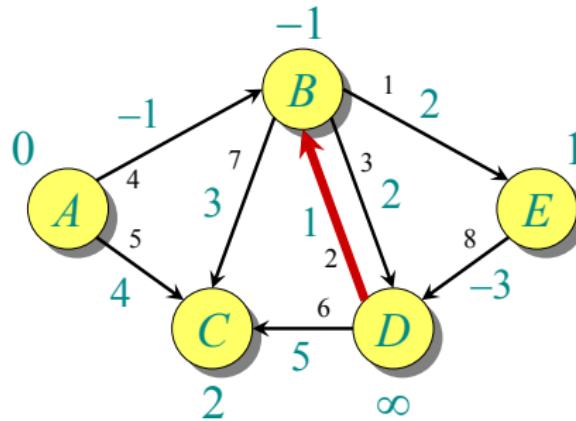


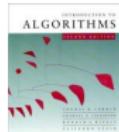
Example of Bellman-Ford



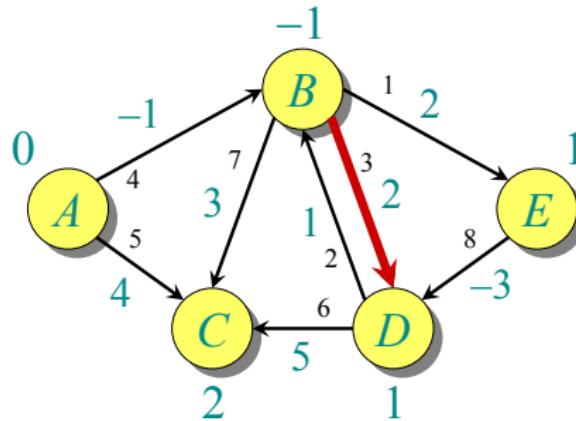


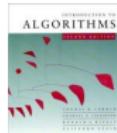
Example of Bellman-Ford



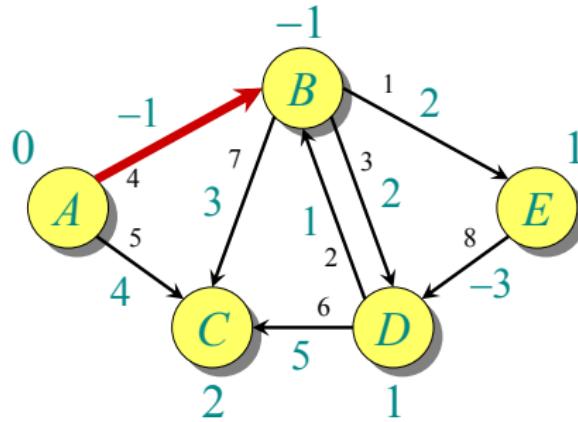


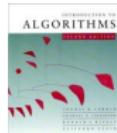
Example of Bellman-Ford



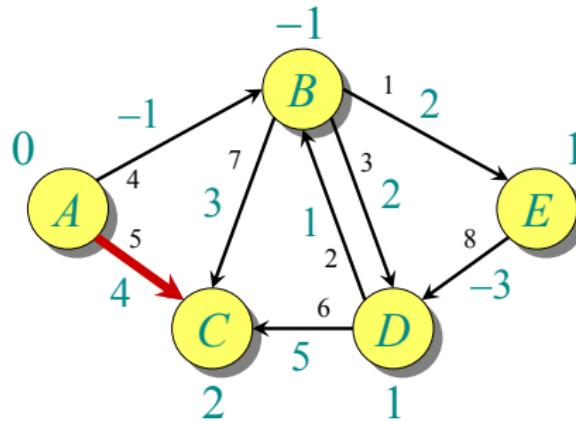


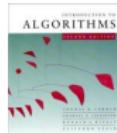
Example of Bellman-Ford



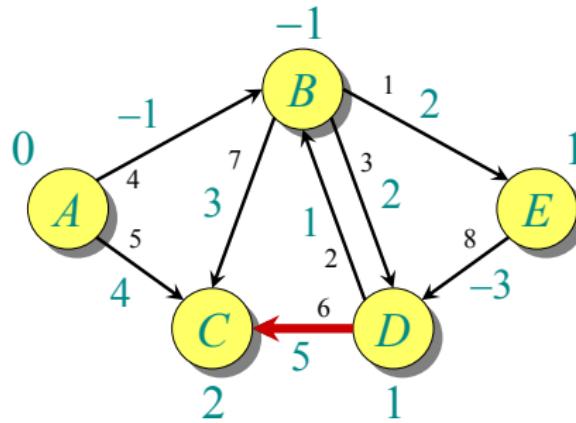


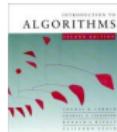
Example of Bellman-Ford



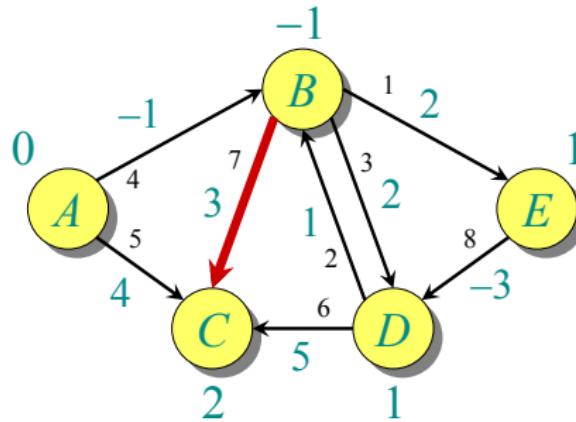


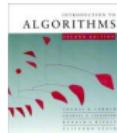
Example of Bellman-Ford



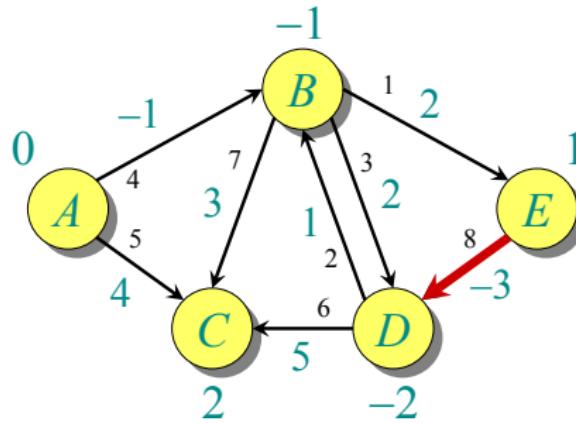


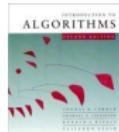
Example of Bellman-Ford



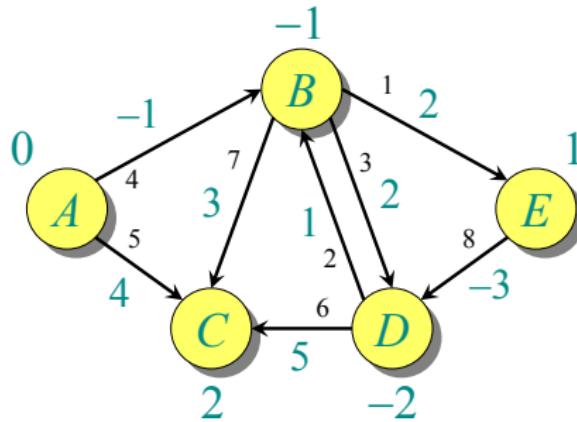


Example of Bellman-Ford





Example of Bellman-Ford



End of pass 2 (and 3 and 4).

Quick question

Is there a more efficient way to find single-source shortest paths for directed acyclic graphs?

Priority first search

- Dijkstra's algorithm for finding single-source shortest paths
- Prim's algorithm for finding minimum spanning trees

are specializations of **priority-first search**.

PRIORITYFIRSTSEARCH(G, s, \dots)

Vertices are visited in order of their **priority**.

For each vertex $v \in G.V$ we calculate:

$v.key$ it's **priority** relative to source vertex s

$v.\pi$ the **predecessor** to v in a **priority-first search tree**

The algorithm uses:

Q a priority queue of unvisited vertices with key $v.key$

Priority-first search

PRIORITYFIRSTSEARCH(G, s, \dots)

- 1 INIT-SINGLE-SOURCE(G, s)
- 2 $Q = G.V$
- 3 **while** $Q \neq \emptyset$
- 4 $u = \text{EXTRACT-MIN}(Q)$
- 5 **for** each vertex $v \in G.Adj[u]$
- 6 RELAX(u, v, \dots)

INIT-SINGLE-SOURCE(G, s)

- 1 **for** each vertex $v \in G.V$
- 2 $v.key = \infty$
- 3 $v.\pi = \text{NIL}$
- 4 $s.key = 0$

RELAX(u, v, \dots)

- 1 **if** $v.key > \text{PRIORITY}$
- 2 $v.key = \text{PRIORITY}$
- 3 $v.\pi = u$

Priority-first search: Prim's algorithm

Instantiate **PRIORITY** = $\text{weight}(u, v)$.

- $G.V - Q$ is the subset of vertices in the MST under construction (T)
- $v.\text{key}$ is the least weight edge connecting v to T
- $v.\pi$ is the vertex adjacent to v on that least-weight edge

At the end, the priority-first search tree is a **minimum spanning tree**.

Priority-first search: Dijkstra's algorithm

Instantiate $\text{PRIORITY} = u.\text{key} + w(u, v)$.

- $G.V - Q$ is the set of vertices already visited in order of their distance from s
- $v.\text{key}$ is the length of the shortest-path from s to v that only uses edges adjacent to visited vertices.
- $v.\pi$ is the predecessor of v on the shortest-path from s to v that only used edges adjacent to visited vertices.

At the end, the priority-first search tree is a **shortest-path-tree**.

Recap of this week

1 Single-source shortest path algorithms:

- Dijkstra's algorithm:
 - For graphs without negative-weight edges
 - $O(E \lg V)$ using binary heap or
 - $O(E + V \lg V)$ (using Fibonacci heap)
- Bellman-ford algorithm:
 - Handles graphs with negative-weight edges
 - Can detect negative-weight cycles
 - $O(VE)$
- Priority-first search:
 - Generalises Prim and Dijkstra's algorithms